

Thesis: A Comparative Study of Large Language Models for Financial Sentiment Analysis and Their Predictive Potential for Short-Term Stock Price Movement

Component: Experiment 1 - Model B : Meta Llama 3 (8B Instruct)

Description: This notebook evaluates Llama 3 8B on the Financial PhraseBank dataset under zero-shot conditions. It generates sentiment predictions and computes accuracy, precision, recall, F1-score (macro-averaged), and confusion matrix.

Select T4 GPU as runtime.

```
# Step 1 - Install bitsandbytes (restart required after this)

!pip install -q -U bitsandbytes accelerate
```

Go to runtime -> restart this session again -> then run cell 1 and run cell 2

```
# Import required libraries - Experiment 1b: Llama 3

import pandas as pd
import numpy as np
import torch
import warnings
import gc

from transformers import AutoTokenizer, AutoModelForCausalLM, BitsAndBytesConfig
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, classification_report
)

warnings.filterwarnings("ignore")

print(f"PyTorch version : {torch.__version__}")
print(f"CUDA available : {torch.cuda.is_available()}")
if torch.cuda.is_available():
    print(f"GPU detected : {torch.cuda.get_device_name(0)}")

PyTorch version : 2.10.0+cu128
CUDA available : True
GPU detected : Tesla T4
```

Add the API Key from Hugging Face using "Add New Secret" in Google Colab

```
# Connect to Hugging Face and load Financial PhraseBank

from google.colab import userdata
from huggingface_hub import login, hf_hub_download

hf_token = userdata.get("HF_TOKEN")
login(token=hf_token)

print("Hugging Face login successful.")

# Load Financial PhraseBank - full dataset
file_path = hf_hub_download(
    repo_id="takala/financial_phrasebank",
    filename="sentences_50agree/train-00000-of-00001.parquet",
    repo_type="dataset",
    revision="0dd3028d70cbd18ded8887e65e83343b03a50482",
    token=hf_token
)

df_fpb = pd.read_parquet(file_path)

label_map = {0: "negative", 1: "neutral", 2: "positive"}
df_fpb["sentiment"] = df_fpb["label"].map(label_map)

true_labels = df_fpb["sentiment"].tolist()
texts = df_fpb["sentence"].tolist()

print(f"\nFinancial PhraseBank loaded.")
print(f"Total sentences : {len(df_fpb)}")
print(f"Label distribution:")
print(df_fpb["sentiment"].value_counts())
```

Hugging Face login successful.

sentences_50agree/train-00000-of-00001.p(...): 100%

390k/390k [00:00<00:00, 1.95MB/s]

Financial PhraseBank loaded.

Total sentences : 4846

Label distribution:

sentiment

neutral 2879

positive 1363

negative 604

Name: count, dtype: int64

```
# Load Llama 3 8B with 4-bit quantisation
```

```
bnb_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_compute_dtype=torch.float16,
    bnb_4bit_use_double_quant=True
)
```

```
print("Loading Llama 3 8B...")
```

```
llama_tokenizer = AutoTokenizer.from_pretrained(
    "meta-llama/Meta-Llama-3-8B-Instruct",
    token=hf_token
)
```

```
llama_model = AutoModelForCausalLM.from_pretrained(
    "meta-llama/Meta-Llama-3-8B-Instruct",
    quantization_config=bnb_config,
    device_map="auto",
    token=hf_token
)
```

```
print("Llama 3 8B loaded successfully.")
```

Loading Llama 3 8B...

config.json: 100%

654/654 [00:00<00:00, 15.8kB/s]

tokenizer_config.json: 51.0k/? [00:00<00:00, 1.24MB/s]

tokenizer.json: 9.09M/? [00:00<00:00, 21.4MB/s]

special_tokens_map.json: 100%

73.0/73.0 [00:00<00:00, 5.72kB/s]

model.safetensors.index.json: 23.9k/? [00:00<00:00, 2.12MB/s]

Download complete: 100% 16.1G/16.1G [04:20<00:00, 92.8MB/s]

Fetching 4 files: 100%

4/4 [04:20<00:00, 260.59s/it]

Loading weights: 100%

291/291 [01:05<00:00, 153.04it/s, Materializing param=model.norm.weight]

generation_config.json: 100%

187/187 [00:00<00:00, 20.2kB/s]

Llama 3 8B loaded successfully.

```
# Define Llama 3 classifier and run on full dataset
```

```
import logging
```

```
logging.getLogger("transformers").setLevel(logging.ERROR)
```

```
def clean_label(label):
    label = label.strip().lower()
    if "positive" in label:
        return "positive"
    elif "negative" in label:
        return "negative"
    elif "neutral" in label:
        return "neutral"
    else:
        return "neutral"
```

```
def classify_llama3(text):
```

```
    messages = [
        {"role": "system", "content": "You are a financial sentiment classifier. Classify the sentiment as exact"},
        {"role": "user", "content": f"Classify the sentiment of this financial text:\n\n{text}\n\nRespond with o"}
    ]
```

```
    tokenized = llama_tokenizer.apply_chat_template(
        messages,
        add_generation_prompt=True,
        return_tensors="pt",
```

```

        return_dict=True
    ).to(llama_model.device)

    with torch.no_grad():
        outputs = llama_model.generate(
            **tokenized,
            max_new_tokens=5,
            do_sample=False,
            pad_token_id=llama_tokenizer.eos_token_id
        )

    input_length = tokenized["input_ids"].shape[1]
    generated = llama_tokenizer.decode(
        outputs[0][input_length:],
        skip_special_tokens=True
    )
    return clean_label(generated)

# Run on full dataset
llama3_preds = []
total = len(texts)

print(f"Running Llama 3 on {total} sentences...")

for i, text in enumerate(texts):
    label = classify_llama3(text)
    llama3_preds.append(label)

    if (i + 1) % 100 == 0:
        print(f" Progress: {i+1}/{total}")

print(f"\nDone. Total predictions: {len(llama3_preds)}")
print(f"\nPrediction distribution:")
print(pd.Series(llama3_preds).value_counts())

```

Running Llama 3 on 4846 sentences...

```

Progress: 100/4846
Progress: 200/4846
Progress: 300/4846
Progress: 400/4846
Progress: 500/4846
Progress: 600/4846
Progress: 700/4846
Progress: 800/4846
Progress: 900/4846
Progress: 1000/4846
Progress: 1100/4846
Progress: 1200/4846
Progress: 1300/4846
Progress: 1400/4846
Progress: 1500/4846
Progress: 1600/4846
Progress: 1700/4846
Progress: 1800/4846
Progress: 1900/4846
Progress: 2000/4846
Progress: 2100/4846
Progress: 2200/4846
Progress: 2300/4846
Progress: 2400/4846
Progress: 2500/4846
Progress: 2600/4846
Progress: 2700/4846
Progress: 2800/4846
Progress: 2900/4846
Progress: 3000/4846
Progress: 3100/4846
Progress: 3200/4846
Progress: 3300/4846
Progress: 3400/4846
Progress: 3500/4846
Progress: 3600/4846
Progress: 3700/4846
Progress: 3800/4846
Progress: 3900/4846
Progress: 4000/4846
Progress: 4100/4846
Progress: 4200/4846
Progress: 4300/4846
Progress: 4400/4846
Progress: 4500/4846
Progress: 4600/4846
Progress: 4700/4846
Progress: 4800/4846

```

Done. Total predictions: 4846

```
Prediction distribution:
neutral      3169
positive     1157
negative      520
Name: count, dtype: int64
```

```
# Evaluate Llama 3 performance on full dataset
```

```
labels_order = ["positive", "negative", "neutral"]
```

```
acc = accuracy_score(true_labels, llama3_preds)
prec = precision_score(true_labels, llama3_preds, average="macro", labels=labels_order)
rec = recall_score(true_labels, llama3_preds, average="macro", labels=labels_order)
f1 = f1_score(true_labels, llama3_preds, average="macro", labels=labels_order)
```

```
print("=" * 45)
print("Llama 3 8B - Experiment 1 Results")
print("=" * 45)
print(f" Accuracy : {acc:.4f} ({acc*100:.2f}%)")
print(f" Precision : {prec:.4f}")
print(f" Recall : {rec:.4f}")
print(f" F1-Score : {f1:.4f}")
print("=" * 45)
```

```
print("\nDetailed Classification Report:")
print(classification_report(true_labels, llama3_preds, labels=labels_order))
```

```
=====
Llama 3 8B - Experiment 1 Results
=====
```

```
Accuracy : 0.7825 (78.25%)
Precision : 0.7867
Recall : 0.7375
F1-Score : 0.7585
=====
```

```
Detailed Classification Report:
```

	precision	recall	f1-score	support
positive	0.74	0.63	0.68	1363
negative	0.83	0.72	0.77	604
neutral	0.79	0.87	0.83	2879
accuracy			0.78	4846
macro avg	0.79	0.74	0.76	4846
weighted avg	0.78	0.78	0.78	4846

```
# Save Llama 3 results
```

```
llama3_scores = {
    "model": "Llama 3 8B",
    "accuracy": round(acc, 4),
    "precision": round(prec, 4),
    "recall": round(rec, 4),
    "f1_score": round(f1, 4)
}
```

```
df_fpb["llama3_pred"] = llama3_preds
```

```
print("Llama 3 8B - Results Summary")
print(pd.DataFrame([llama3_scores]))
```

```
Llama 3 8B - Results Summary
```

	model	accuracy	precision	recall	f1_score
0	Llama 3 8B	0.7825	0.7867	0.7375	0.7585

```
# Save Llama 3 predictions to Google Drive
```

```
from google.colab import drive
drive.mount("/content/drive")
```

```
import os
os.makedirs("/content/drive/MyDrive/Thesis_Data", exist_ok=True)
```

```
df_fpb.to_csv("/content/drive/MyDrive/Thesis_Data/exp1_llama3_preds.csv", index=False)
```

```
print("Llama 3 predictions saved to Google Drive.")
```

```
Mounted at /content/drive
Llama 3 predictions saved to Google Drive.
```


```
# Confusion Matrix – Llama 3 8B Experiment 1
```

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

labels_order = ["positive", "negative", "neutral"]

cm = confusion_matrix(true_labels, llama3_preds, labels=labels_order)

plt.figure(figsize=(7, 5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=labels_order,
            yticklabels=labels_order)
plt.xlabel("Predicted Label", fontsize=12)
plt.ylabel("True Label", fontsize=12)
plt.title("Llama 3 8B – Confusion Matrix (Experiment 1)", fontsize=12)
plt.tight_layout()
plt.savefig("/content/drive/MyDrive/Thesis_Data/fig_cm_llama3.png",
            dpi=300, bbox_inches="tight")
plt.show()
print("Confusion matrix saved to Google Drive.")
```



Confusion matrix saved to Google Drive.